

Movie Review Analysis using PIG

Project submitted to the
SRM University – AP, Andhra Pradesh
for the partial fulfillment of the requirements to award the degree of

Bachelor of Technology

In

**Computer Science and Engineering
School of Engineering and Sciences**

Submitted by

G. Mayur Teja (AP20110010087)

G. Asish Karthikeya (AP20110010114)

K. Harika (AP20110010124)



Under the Guidance of

[Sriramulu Bojjagani](#)

SRM University-AP

Neerukonda, Mangalagiri, Guntur

Andhra Pradesh – 522 240

[November, 2023]

Certificate

Date: 30-Oct-23

This is to certify that the work present in this Project entitled “**Movie Review Analysis Using PIG**” has been carried out by **G. Mayur teja, G. Asish Karthikeya, K. Harika** under my supervision. The work is genuine, original, and suitable for submission to the SRM University – AP for the award of Bachelor of Technology in the **School of Engineering and Sciences**.

Supervisor

(Signature)

Dr. Sriramulu Bojjagani

Assistant Professor,

Computer Science and Engineering.

Acknowledgements

I am grateful to Dr.Sriramulu Bojjagani for his ongoing support and fascinating report techniques. I also want to express my appreciation for them giving us their pearls of wisdom. I am grateful for his willingness to give up his time so freely.

I would also like to express my gratitude to SRM University, AP for assistance with analyzing the meteorological data, and for providing a good environment and facilities to complete this report successfully. I also appreciate the thoughtful criticism from the Books & Texts peer reviewers who remain anonymous.

I appreciate having the opportunity to work on this with everyone. My dissertation committee members have all given me a great deal of professional and personal advice and have taught me a lot about both scientific research and life in general.

Finally, I want to thank my parents for their inspiration and support throughout my academic career.

Table of Contents

Certificate	1
Acknowledgements	2
Table of Contents	3
Abstract	4
Abbreviations	5
List of Figures	6
1.Introduction	7
1.1 Applications:	8
1.2 PIG: Working and Environment.	9
2.Methodology	11
2.1 Dataset	11
2.2.1 Movie Analysis	12
2.2.2 Grouping and Aggregating: You can group the data by certain fields and perform aggregate operations. Finding the average rating for each user.	17
3. Concluding Remark	24
4. Future Work	25
5. References	26

Abstract

Movies have constantly functioned as a key source of entertainment since its debut in the late 18th century, delivering a diverse range of genres such as drama, comedy, science fiction, and action. The vast and diversified nature of acquired movie data needs a departure from standard analytical methodologies, advocating the use of big data analytics for complete insights. This research digs into the investigation and analysis of a movie ratings dataset using HDFS and Apache Pig, a platform for large-scale data processing. Using Pig's data processing capability, the research focuses on identifying movies with low average ratings, particularly those below 2.0, also Additionally, arranging the top 10 movies based on their highest average ratings. and gives a comprehensive assessment of their titles, average ratings, and number of ratings received. The findings shed light on cinematic works that may not have connected well with audiences, providing vital insights into the dynamics of movie tastes, and contributing to the expanding environment of data-driven film study.

Keywords: Big data, HDFS, Apache PIG

Abbreviations

SQL	Structured Query Language
HDFS	Hadoop Distributed File System

List of Figures

Figure 1. Research Framework	08
Figure 2. Sandbox Login.....	12
Figure 3. Loading Dataset.....	13
Figure 4 Query.....	13
Figure 5. Storing the result.....	14
Figure 6. Query.....	18
Figure 7. Query.....	20
Figure 8. Query.....	21
Figure 9. Query.....	22

1.Introduction

The growth of user-generated material in the digital age has altered the landscape of information sharing and entertainment consumption. The massive volume of data created by internet movie reviews is one visible indication of this transition. Moviegoers share their thoughts on numerous platforms, producing a wealth of useful information for filmmakers, studios, and fans. It is essential to use big data technology to analyze movie reviews to get valuable insights from this enormous and dynamic dataset. The expansion of digital platforms has resulted in a flood of user-generated material, particularly in the form of movie reviews. These evaluations, which can be found on a variety of websites and forums, include a lot of information that may be used to identify patterns, trends, and attitudes associated to cinematic experiences. When confronted with the vast volume of data accessible, traditional techniques of analysis fall short, necessitating the use of modern big data technology.

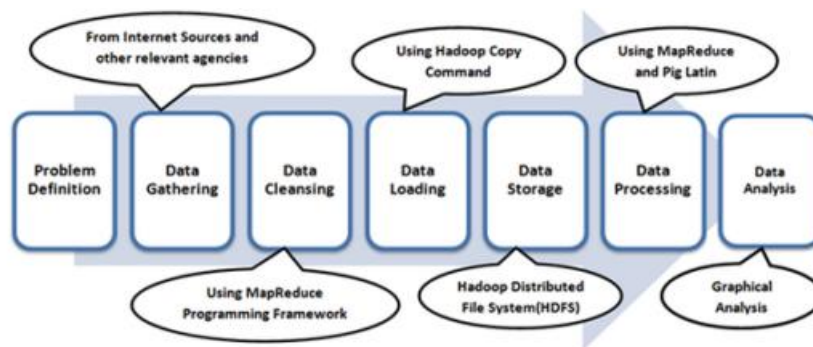
Our study focuses on the use of Apache Pig and MapReduce Hadoop, two strong big data tools, to process and evaluate massive repositories of movie reviews. Apache Pig provides a high-level framework for expressing data analysis applications, allowing researchers to create complicated analysis workflows without wading into MapReduce programming complexities. Hadoop, on the other hand, is a distributed processing platform that is well-known for its scalability and efficiency when dealing with enormous datasets. The convergence of these technologies allows us to address the issues posed by large datasets by providing a scalable and parallelized method to movie review research. Our technique entails extracting important characteristics from textual data, sentiment analysis, and identifying major themes within reviews. We want to deliver a complete and effective solution for extracting meaningful insights from the large ocean of movie-related thoughts by using the parallel processing capabilities of MapReduce Hadoop and the data manipulation ease of Apache Pig.

In the subsequent sections, we will delve into the intricacies of our methodology, present the results of our analyses, and discuss the implications and potential applications of our findings in the broader context of the film industry. Through this research, we endeavor to demonstrate the efficacy of employing Pig and MapReduce Hadoop in the domain of movie review analysis, opening new avenues for the exploration of audience sentiments in the digital age.

1.1 Applications:

- **Movie Recommendation System:** Analyzing movie ratings and identifying poorly rated movies can be a crucial step in enhancing a movie recommendation system. The application could use this analysis to filter out poorly rated movies from recommendations to improve user satisfaction.
- **Content Quality Assessment:** Media companies could use this script to assess the overall quality of their movie catalog. Identifying movies with low average ratings can help in deciding whether to promote, demote, or remove certain content from their offerings.
- **User Engagement Analysis:** For streaming platforms, understanding the average ratings and the number of ratings per movie can provide insights into user engagement. Low-rated movies with a high number of ratings may indicate controversial content that attracts audience attention.

Figure 1. Research Framework:



1.2 PIG: Working and Environment.

Apache Pig serves as an elevated platform crafted for the development of MapReduce programs within the Hadoop ecosystem, streamlining the intricate processes associated with large-scale data processing by offering an abstraction layer over the intricacies of traditional MapReduce programming. At its core lies Pig Latin, a scripting language meticulously designed to articulate data transformations in a concise and adaptable manner.

1. Pig Latin and Data Transformations:

Pig Latin Language: Functioning as a procedural scripting language, Pig Latin facilitates the articulation of data transformations. Its design prioritizes intuition over the complexities inherent in traditional MapReduce, enabling developers to concentrate on the logic of data processing rather than the nuances of parallelization and optimization.

Data Transformation Tasks: Pig scripts find particular efficacy in tasks such as Extract, Transform, Load (ETL), data analysis, and iterative data processing. Pig accommodates various data types, encompassing both structured and unstructured data, while providing an array of operations including filtering, joining, ordering, and more.

2. Data Sources and Versatility:

Data Sources: Pig seamlessly manages data from diverse sources, spanning databases, local files, and Hadoop's Distributed File System (HDFS). This adaptability positions it as a valuable tool for organizations dealing with a myriad of data formats and storage systems.

Extensibility: Pig's functionality can be extended by developers through the creation of their own user-defined functions (UDFs) for tasks such as reading, processing, and writing data. This inherent extensibility enhances Pig's adaptability to a spectrum of data processing requirements.

3. Parallel Execution and Optimization:

Parallel Processing: A standout feature of Pig is its adeptness in handling parallel processing. Tasks are automatically optimized by breaking them into sub-tasks that can execute concurrently across a Hadoop cluster. This capability is indispensable for managing voluminous datasets and leveraging the distributed architecture of Hadoop.

Task Optimization: Pig abstracts away the intricacies associated with parallel execution, empowering developers to devise complex data transformations without delving into the underlying architectural complexities or optimization intricacies. This characteristic renders Pig an accessible tool for a diverse range of data scientists and engineers.

4. Use Cases and Popularity:

ETL and Data Analysis: Pig finds common application in ETL tasks, where data is extracted from diverse sources, transformed, and loaded into a data warehouse or analytical system. Its straightforward nature also makes it a preferred choice for exploratory data analysis.

Iterative Data Processing: The platform's support for iterative data processing renders it suitable for algorithms requiring multiple iterations, such as machine learning algorithms.

Ease of Use: Pig's simplicity and flexibility, coupled with its support for parallel processing, position it as a favored choice among data scientists and engineers working with substantial datasets. This allows them to emphasize data analysis while minimizing the complexities associated with writing MapReduce programs.

In summary, Apache Pig's capacity to handle diverse data types, facilitate complex data operations, and execute tasks in parallel establishes it as an indispensable tool within the Hadoop ecosystem. Its design principles of simplicity, extensibility, and parallelism contribute significantly to its popularity among developers and data practitioners involved in large-scale

2.Methodology

2.1 Dataset

This research leverages two distinct datasets, comprises movie-related information and is sourced from “**github**”. The data includes details on various movies, such as titles, release dates, IMDb URLs, and genre classifications.

1. Movie Metadata Dataset

The first dataset, we referred it as the "Movie Metadata Dataset, It includes comprehensive information about various movies, such as titles, release dates, IMDb URLs, and genre classifications.

The dataset is structured in tabular form, with each row representing a distinct movie. The columns include:

MovieID: A unique identifier for each movie.

Title: The title of the movie.

ReleaseDate: The date of the movie's release.

URL: The IMDb URL associated with the movie.

Genres: Multiple genre columns (Genre1, Genre2, ..., GenreN) indicating the genres of each movie.

2. User Ratings Dataset

The second dataset we referred it as "User Ratings Dataset,"

Similar to the Movie Metadata Dataset, the User Ratings Dataset is structured in tabular form. The columns include:

UserID: A unique identifier for each user.

MovieID: The unique identifier corresponding to each movie.

Rating: The user-assigned rating for the movie on a numerical scale.

2.2 Procedure

2.2.1 Movie Analysis

1. **Data Ingestion and HDFS Setup:** Begin by setting up the Hadoop Distributed File System (HDFS) environment. Upload the provided datasets mentioned earlier into the newly created folder using the 'get' command.

```
sandbox login: root
root@sandbox.hortonworks.com's password:
Last login: Tue Sep 19 02:09:21 2023 from 172.17.0.2
[root@sandbox ~]# hadoop fs -get /Project_Dataset/u.item
get: `u.item': File exists
[root@sandbox ~]# hadoop fs -get /Project_Dataset/u.item
get: `u.item': File exists
[root@sandbox ~]#
```

Figure 2. Sandbox Login & Data Ingestion

2. **Launching Pig:** Open the PIG shell by typing 'pig' in the command line interface.

```
[root@sandbox ~]# pig
23/09/19 02:13:49 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL
23/09/19 02:13:49 INFO pig.ExecTypeProvider: Trying ExecType : MAPREDUCE
23/09/19 02:13:49 INFO pig.ExecTypeProvider: Picked MAPREDUCE as the ExecType
2023-09-19 02:13:50,047 [main] INFO org.apache.pig.Main - Apache Pig version 0.16.0.2.5.0.0-1245 (rexported) compiled Aug 26 2016, 02:07:35
2023-09-19 02:13:50,048 [main] INFO org.apache.pig.Main - Logging error messages to: /root/pig_1695089630045.log
2023-09-19 02:13:50,079 [main] INFO org.apache.pig.impl.util.Utils - Default bootstrap file /root/.pigbootstrap not found
2023-09-19 02:13:50,708 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to hadoop file system at: hdfs://sandbox.hortonworks.com:8020
2023-09-19 02:13:51,339 [main] INFO org.apache.pig.PigServer - Pig Script ID for the session: PIG-default-1c846bf0-c7c3-49fd-bda9-a4db277ba15e
2023-09-19 02:13:51,674 [main] INFO org.apache.hadoop.yarn.client.api.impl.TimelineClientImpl - Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
2023-09-19 02:13:51,791 [main] INFO org.apache.pig.backend.hadoop.PigATSCClient - Created ATS Hook
grunt>
```

3. **Loading the Datasets:** The first line of the script loads the data located at 'Project_Dataset'. The load function is used to load the data and PigStorage is used to specify the delimiter (comma in this case).

```
grunt> ratings = LOAD '/Project_Dataset/u.data' AS (userID:int, movieID:int, rating:
int, ratingTime:int);
grunt> metadata = LOAD '/Project_Dataset/u.item' USING PigStorage(',') AS (movieID:
int, movieTitle:chararray, releaseDate:chararray, videoRelease:chararray, imdbLink:
chararray);
```

Figure 3 .Loading Dataset

4. **Projection Operation (nameLookup):** In this step, a new relation named nameLookup is created by selecting only the movieID and movieTitle columns from the metadata relation.
5. **Grouping Operation (groupedRatings):** The ratings relation is grouped by the movieID field. This is a typical step before performing operations on groups of data.
6. **Average Calculation (avgRatings):** This step calculates the average rating (avgRating) and the total number of ratings (numRatings) for each movie. The results are stored in the avgRatings relation.
7. **Filtering Operation (badMovies):** Movies with an average rating less than 2.0 are filtered and stored in the badMovies relation.
8. **Join Operation (namedBadMovies):** The badMovies relation is joined with the nameLookup relation based on the common field movieID.
9. **Projection Operation (finalResults):** In this step, unnecessary columns are removed, and only the required columns (movieName, avgRating, numRatings) are selected for the final results.
10. **Ordering Operation (finalResultsSorted):** The final results are ordered in descending order based on the number of ratings (numRatings).

```
grunt> nameLookup = FOREACH metadata GENERATE movieID, movieTitle;
grunt> groupedRatings = GROUP ratings BY movieID;
grunt> avgRatings = FOREACH groupedRatings GENERATE group AS movieID, AVG(ratings.r
ating) AS avgRating, COUNT(ratings.rating) AS numRatings;
grunt> badMovies = FILTER avgRatings BY avgRating < 2.0;
grunt> namedBadMovies = JOIN badMovies BY movieID, nameLookup BY movieID;
grunt> finalResults = FOREACH namedBadMovies GENERATE nameLookup::movieTitle AS mov
ieName, badMovies::avgRating AS avgRating,
>>
>> badMovies::numRatings AS numRatings;
grunt> finalResultsSorted = ORDER finalResults BY numRatings DESC;
```

Figure 4.Query

12. Storing the Results: The results are stored in to separate text file located at '/Project_Dataset/finalResultsSorted.txt'. The 'Store' function is used to store the results.

```
grunt> STORE finalResultsSorted INTO '/Project_Dataset/finalResultsSorted.txt';
2023-09-19 01:27:30,006 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: HASH_JOIN,GROUP_BY,ORDER_BY,FILTER
2023-09-19 01:27:30,053 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2023-09-19 01:27:30,138 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, ConstantCalculator, GroupByConstParallelSetter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, PartitionFilterOptimizer, PredicatePushdownOptimizer, PushDownForEachFilter, PushUpFilter, SplitFilter, StreamTypeCastInserter]}
2023-09-19 01:27:30,189 [main] INFO org.apache.pig.newplan.logical.rules.ColumnPruneVisitor - Columns pruned for metadata: $2, $3, $4
2023-09-19 01:27:30,284 [main] INFO org.apache.pig.impl.util.SpillableMemoryManager - Selected heap (PS Old Gen) of size 699400192 to monitor. collectionUsageThreshold = 489580128, usageThreshold = 489580128
2023-09-19 01:27:30,375 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2023-09-19 01:27:30,426 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.CombinerOptimizerUtil - Choosing to move algebraic foreach to combiner
2023-09-19 01:27:30,446 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.SecondaryKeyOptimizerMR - Using SecondaryKey Optimization for MapReduce node scope-71
2023-09-19 01:27:30,452 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler$LastInputStreamingOptimizer - Rewrite: POPackage->POForEach to POPackage(JoinPackager)
2023-09-19 01:27:30,465 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 4
2023-09-19 01:27:30,466 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 4
2023-09-19 01:27:30,600 [main] INFO org.apache.hadoop.yarn.client.api.impl.TimelineClientImpl - Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
2023-09-19 01:27:30,619 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at sandbox.hortonworks.com/172.17.0.2:8050
2023-09-19 01:27:30,812 [main] INFO org.apache.hadoop.yarn.client.AHSPProxy - Connecting to Application History server at sandbox.hortonworks.com/172.17.0.2:10200
2023-09-19 01:27:30,912 [main] INFO org.apache.pig.tools.pigstats.mapreduce.MRScriptState - Pig script settings are added to the job
2023-09-19 01:27:30,922 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2023-09-19 01:27:30,928 [main] INFO org.apache.pig.backend.hadoop.executionengine.
```

Figure 5. Storing the result

Output of this query :

```
Records Written : 100000
HadoopVersion  PigVersion  UserId  StartedAt  FinishedAt  Features
2.7.3.2.5.0.0-1245  0.16.0.2.5.0.0-1245  root  2023-09-19 01:47:49  2023-09-19 01:48:17  HASH_JOIN
Success!

Job Stats (time in seconds):
JobId  Maps  Reduces  MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime  MaxReduceTime  MinReduceTime  AvgReduceTime  MedianReducetime  Alias  F
eature  Outputs
job_1694535283099_0062  2  1  5  5  5  5  5  5  5  5  metadata,ratings,ratings_with_titles  HASH_JOIN  /Project
_Dataset/ratings_with_titles.txt,

Input(s):
Successfully read 1682 records from: "/Project_Dataset/u.item"
Successfully read 100000 records from: "/Project_Dataset/u.data"

Output(s):
Successfully stored 100000 records (12163405 bytes) in: "/Project_Dataset/ratings_with_titles.txt"

Counters:
Total records written : 100000
Total bytes written : 12163405
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1694535283099_0062
```

File Preview ✕

/Project_Dataset/finalResultsSorted.txt/part-r-00000

Leave It to Beaver (1997)	1.8409090909090908	44	
Mortal Kombat: Annihilation (1997)	1.9534883720930232	43	
Crow: City of Angels, The (1996)	1.9487179487179487	39	
Bio-Dome (1996)	1.903225806451613	31	
Barb Wire (1996)	1.9333333333333333	30	
Free Willy 3: The Rescue (1997)	1.7407407407407407	27	
Showgirls (1995)	1.9565217391304348	23	
Lawnmower Man 2: Beyond Cyberspace (1996)	1.7142857142857142	21	
Children of the Corn: The Gathering (1996)	1.3157894736842106	19	
Home Alone 3 (1997)	1.894736842105263	19	
Ready to Wear (Pret-A-Porter) (1994)	1.8333333333333333	18	
Jaws 3-D (1983)	1.9375	16	
All Dogs Go to Heaven 2 (1996)	1.8666666666666667	15	
Big Bully (1996)	1.8571428571428572	14	
Amityville II: The Possession (1982)	1.6428571428571428	14	
Body Parts (1991)	1.6153846153846154	13	
Mr. Magoo (1997)	1.9166666666666667	12	

✕ Cancel

Download

(Leave It to Beaver (1997),1.8409090909090908,44)
 (Mortal Kombat: Annihilation (1997),1.9534883720930232,43)
 (Crow: City of Angels, The (1996),1.9487179487179487,39)
 (Bio-Dome (1996),1.903225806451613,31)
 (Barb Wire (1996),1.9333333333333333,30)
 (Free Willy 3: The Rescue (1997),1.7407407407407407,27)
 (Showgirls (1995),1.9565217391304348,23)
 (Lawnmower Man 2: Beyond Cyberspace (1996),1.7142857142857142,21)
 (Children of the Corn: The Gathering (1996),1.3157894736842106,19)
 (Home Alone 3 (1997),1.894736842105263,19)
 (Ready to Wear (Pret-A-Porter) (1994),1.8333333333333333,18)
 (Jaws 3-D (1983),1.9375,16)
 (All Dogs Go to Heaven 2 (1996),1.8666666666666667,15)
 (Big Bully (1996),1.8571428571428572,14)
 (Amityville II: The Possession (1982),1.6428571428571428,14)
 (Body Parts (1991),1.6153846153846154,13)
 (Mr. Magoo (1997),1.9166666666666667,12)
 (Solo (1996),1.8333333333333333,12)
 (Vampire in Brooklyn (1995),1.8333333333333333,12)
 (Robocop 3 (1993),1.7272727272727273,11)
 (Gone Fishin' (1997),1.8181818181818181,11)
 (Kazaam (1996),1.8,10)
 (Bloodsport 2 (1995),1.7,10)
 (Pest, The (1997),1.875,8)
 (Dangerous Ground (1997),1.75,8)
 (Dunston Checks In (1996),1.8571428571428572,7)
 (Bushwhacked (1995),1.5714285714285714,7)
 (Ed (1996),1.3333333333333333,6)
 (Land Before Time III: The Time of the Great Giving (1995) (V),1.6666666666666667,6)
 (Best of the Best 3: No Turning Back (1995),1.5,6)
 (Amityville 3-D (1983),1.1666666666666667,6)
 (Bad Girls (1994),1.8333333333333333,6)
 (Turbo: A Power Rangers Movie (1997),1.8,5)
 (3 Ninjas: High Noon At Mega Mountain (1998),1.0,5)
 (Theodore Rex (1995),1.4,5)
 (Amityville 1992: It's About Time (1992),1.0,5)
 (Phat Beach (1996),1.6,5)
 (Amityville: A New Generation (1993),1.0,5)
 (Talking About Sex (1994),1.8,5)
 (Getting Even with Dad (1994),1.6,5)
 (Amityville Curse, The (1990),1.25,4)
 (Designated Mourner, The (1997),1.75,4)

```

(Phat Beach (1996),1.6,5)
(Amityville: A New Generation (1993),1.0,5)
(Talking About Sex (1994),1.8,5)
(Getting Even with Dad (1994),1.6,5)
(Amityville Curse, The (1990),1.25,4)
(Designated Mourner, The (1997),1.75,4)
(Castle Freak (1995),1.25,4)
(Neon Bible, The (1995),1.75,4)
(Being Human (1993),1.75,4)
(Relative Fear (1994),1.6666666666666667,3)
(Gordy (1995),1.0,3)
(Squeeze (1996),1.6666666666666667,3)
(Catwalk (1995),1.3333333333333333,3)
(Temptress Moon (Feng Yue) (1996),1.6666666666666667,3)
(Simple Wish, A (1997),1.6666666666666667,3)
(Amityville: Dollhouse (1996),1.0,3)
(B*A*P*S (1997),1.6666666666666667,3)
(Mighty, The (1998),1.0,3)
(Gilligan's Island: The Movie (1998),1.3333333333333333,3)
(Time Tracers (1995),1.5,2)
(Country Life (1994),1.5,2)
(Carmen Miranda: Bananas Is My Business (1994),1.5,2)
(Naked in New York (1994),1.5,2)
(Sexual Life of the Belgians, The (1994),1.5,2)
(Boys in Venice (1996),1.0,2)
(Venice/Venice (1992),1.0,2)
(Somebody to Love (1994),1.0,2)
(Crows and Sparrows (1949),1.5,2)
(Babyfever (1994),1.0,2)
(Dream Man (1995),1.5,2)
(Sudden Manhattan (1996),1.5,2)
(Falling in Love Again (1980),1.0,2)
(Johnny 100 Pesos (1993),1.5,2)
(Beyond Bedlam (1993),1.0,2)
(Machine, The (1994),1.5,2)
(Mat' i syn (1997),1.0,1)
(Further Gesture, A (1996),1.0,1)
(New Age, The (1994),1.0,1)
(Getting Away With Murder (1996),1.0,1)
(Chairman of the Board (1998),1.0,1)
(Nobody Loves Me (Keiner liebt mich) (1994),1.0,1)
(Butterfly Kiss (1995),1.0,1)

```

Output of the query with movie rating <2.0

2.2.2 Grouping and Aggregating: You can group the data by certain fields and perform aggregate operations. Finding the average rating for each user.

Script:

```
ratings = LOAD '/Project_Dataset/u.data' AS (userID:int, movieID:int, rating:int, ratingTime:int);
```

```
grouped_ratings_by_user = GROUP ratings BY userID;
```

```
avg_ratings_by_user = FOREACH grouped_ratings_by_user GENERATE group AS userID, AVG(ratings.rating) AS avgRating;
```

```
STORE avg_ratings_by_user INTO '/Project_Dataset/avg_ratings_by_user.txt';
```

```
grunt> ratings = LOAD '/Project_Dataset/u.data' AS (userID:int, movieID:int, rating
:int, ratingTime:int);
grunt> grouped_ratings_by_user = GROUP ratings BY userID;
grunt> avg_ratings_by_user = FOREACH grouped_ratings_by_user GENERATE group AS user
ID, AVG(ratings.rating) AS avgRating;
grunt> STORE avg_ratings_by_user INTO '/Project_Dataset/avg_ratings_by_user.txt';
```

Figure 6 query

Output:

```
HadoopVersion  PigVersion      UserId StartedAt      FinishedAt      Features
2.7.3.2.5.0.0-1245  0.16.0.2.5.0.0-1245  root  2023-09-19 02:50:27  202
3-09-19 02:51:05      GROUP_BY

Success!

Job Stats (time in seconds):
JobId  Maps  Reduces MaxMapTime  MinMapTime  AvgMapTime  MedianMapTi
me      MaxReduceTime  MinReduceTime  AvgReduceTime  MedianReductime  Ali
as      Feature Outputs
job_1694535283099_0071  1      1      6      6      6      6      11      111
1      11      avg_ratings_by_user,grouped_ratings_by_user,ratings  GROUP_BY,CO
MBINER  /Project_Dataset/avg_ratings_by_user.txt,

Input(s):
Successfully read 100000 records (1979558 bytes) from: "/Project_Dataset/u.data"

Output(s):
Successfully stored 943 records (19336 bytes) in: "/Project_Dataset/avg_ratings_by_
user.txt"

Counters:
Total records written : 943
Total bytes written : 19336
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1694535283099_0071
```

File Preview

/Project_Dataset/avg_ratings_by_user.txt/part-r-00000

1	3.610294117647059
2	3.7096774193548385
3	2.7962962962962963
4	4.333333333333333
5	2.874285714285714
6	3.6350710900473935
7	3.965260545905707
8	3.7966101694915255
9	4.2727272727272725
10	4.206521739130435
11	3.4640883977900554
12	4.392156862745098
13	3.09748427672956
14	4.091836734693878
15	2.875
16	4.328571428571428
17	3.0357142857142856
18	3.88086642599278

2.2.3 Joining: You can join multiple datasets together. Joining the ratings and metadata to get movie titles in the ratings.

Script:

ratings_with_titles = JOIN ratings BY movieID, metadata BY movieID;

STORE ratings_with_titles INTO '/Project_Dataset/ratings_with_titles.txt';

```
grunt> ratings_with_titles = JOIN ratings BY movieID, metadata BY movieID;
grunt> STORE ratings_with_titles INTO '/Project_Dataset/ratings_with_titles.txt';
2023-09-19 01:47:49,512 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: HASH_JOIN
2023-09-19 01:47:49,544 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2023-09-19 01:47:49,544 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, ConstantCalculator, GroupByConstParallelSetter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, PartitionFilterOptimizer, PredicatePushdownOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter]}
2023-09-19 01:47:49,545 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2023-09-19 01:47:49,546 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler$LastInputStreamingOptimizer - Rewrite: POPackage->POForEach to POPackage(JoinPacker)
2023-09-19 01:47:49,546 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2023-09-19 01:47:49,546 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2023-09-19 01:47:49,629 [main] INFO org.apache.hadoop.yarn.client.api.impl.TimelineClientImpl - Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
2023-09-19 01:47:49,629 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at sandbox.hortonworks.com/172.17.0.2:8050
2023-09-19 01:47:49,630 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to Application History server at sandbox.hortonworks.com/172.17.0.2:10200
2023-09-19 01:47:49,637 [main] INFO org.apache.pig.tools.pigstats.mapreduce.MRScriptState - Pig script settings are added to the job
2023-09-19 01:47:49,638 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2023-09-19 01:47:49,638 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Reduce phase detected, estimating # of required reducers.
2023-09-19 01:47:49,638 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator
2023-09-19 01:47:49,642 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator - BytesPerReducer=1000000000 maxReducers=999 totalInputFileSize=2215517
2023-09-19 01:47:49,642 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting Parallelism to 1
2023-09-19 01:47:49,642 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - This job cannot be converted run in-process
2023-09-19 01:47:49,774 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Added jar file:/usr/hdp/2.5.0.0-1245/pig/pig-0.16.0.2.5.0.0-1245-core-h2.jar to DistributedCache through /tmp/temp-1280146786/tmp2100227063/pig-0.16.0.2.5.0.0-1245-core-h2.jar
2023-09-19 01:47:49,832 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Added jar file:/usr/hdp/2.5.0.0-1245/pig/lib/automaton-1.11-8.jar to DistributedCache through /tmp/temp-1280146786/tmp291009866/automaton-1.11-8.jar
2023-09-19 01:47:49,916 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Added jar file:/usr/hdp/2.5.0.0-1245/pig/lib/antlr-runtime-3.4.jar to DistributedCache through /tmp/temp-1280146786/tmp1701875238/antlr-runtime-3.4.jar
2023-09-19 01:47:50,408 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Added jar file:/usr/hdp/2.5.0.0-1245/hadoop/lib/joda-time-2.8.1.jar to DistributedCache through /tmp/temp-1280146786/tmp-969661454/joda-time-2.8.1.jar
2023-09-19 01:47:50,401 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2023-09-19 01:47:50,401 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2023-09-19 01:47:50,401 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2023-09-19 01:47:50,401 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2023-09-19 01:47:50,409 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2023-09-19 01:47:50,467 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.TimelineClientImpl - Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
```

Figure 7 Query

Output:

```
HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.7.3.2.5.0.0-1245 0.16.0.2.5.0.0-1245 root 2023-09-19 01:47:49 2023-09-19 01:48:17 HASH_JOIN

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReductime Alias F
eature Outputs
job_1694535283099_0062 2 1 5 5 5 5 5 5 5 metadata,ratings,ratings_with_titles HASH_JOIN /Project_Dataset/ratings_with_titles.txt,

Input(s):
Successfully read 1682 records from: "/Project_Dataset/u.item"
Successfully read 100000 records from: "/Project_Dataset/u.data"

Output(s):
Successfully stored 100000 records (12163405 bytes) in: "/Project_Dataset/ratings_with_titles.txt"

Counters:
Total records written : 100000
Total bytes written : 12163405
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1694535283099_0062
```

892	1	5	886608185	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
622	1	3	882590344	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
247	1	4	893097024	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
815	1	5	878691975	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
779	1	4	875501555	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
184	1	4	889907652	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
15	1	1	879455635	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
338	1	3	879438143	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
883	1	3	891914583	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
471	1	4	889827881	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
286	1	4	876521699	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
320	1	3	884748604	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
120	1	4	889490412	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
612	1	4	875324876	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
177	1	3	880130699	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
680	1	4	876816224	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti
807	1	4	892528231	1	Toy Story (1995)	01-Jan-1995	http://us.imdb.com/M/ti

2.2.4 Sorting: You can sort the data based on certain fields. Sorting the movies based on their average ratings.

Script:

sorted_movies = ORDER avg_ratings_by_user BY avgRating DESC;

STORE sorted_movies INTO '/Project_Dataset/sorted_movies.txt';

```
grunt> sorted_movies = ORDER avg_ratings_by_user BY avgRating DESC;
grunt> STORE sorted_movies INTO '/Project_Dataset/sorted_movies.txt';
2023-09-19 03:05:31,777 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig
features used in the script: GROUP_BY,ORDER_BY
2023-09-19 03:05:31,835 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [
```

Figure 8 Query

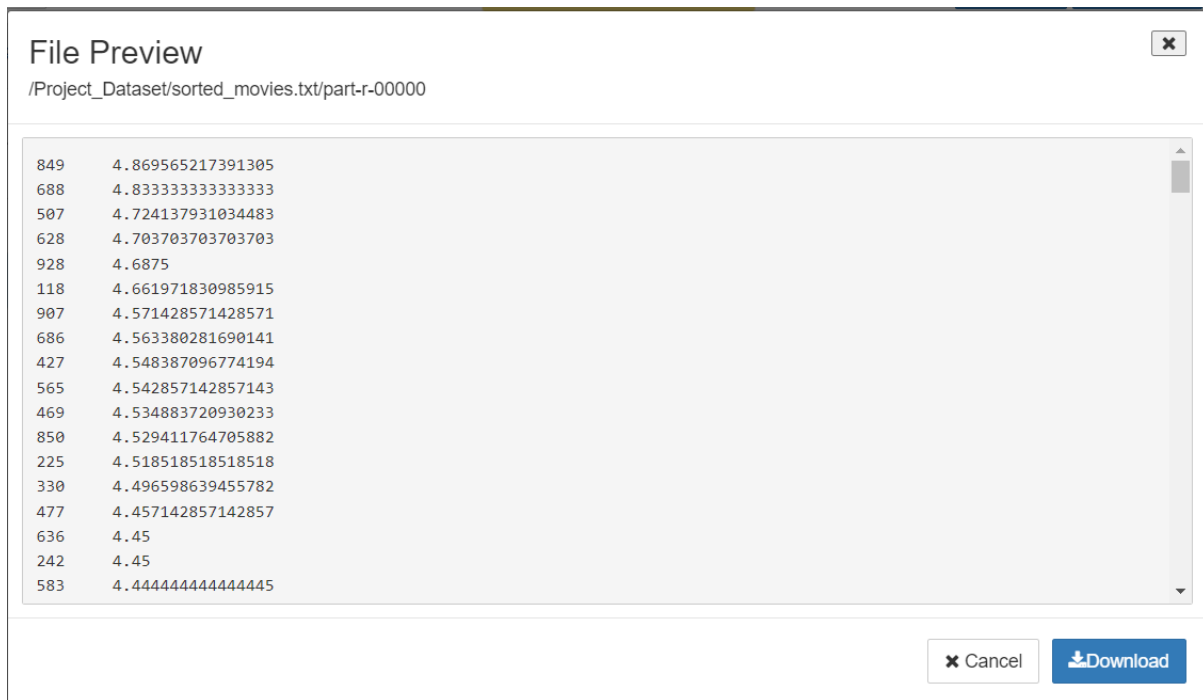
Output:

```
Input(s):
Successfully read 100000 records (1979558 bytes) from: "/Project_Dataset/u.data"

Output(s):
Successfully stored 943 records (19336 bytes) in: "/Project_Dataset/sorted_movies.txt"

Counters:
Total records written : 943
Total bytes written : 19336
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1694535283099_0075 -> job_1694535283099_0076,
job_1694535283099_0076 -> job_1694535283099_0077,
job_1694535283099_0077
```



2.2.5 Limiting: You can limit the number of tuples in a relation. Can get the top 10 movies with the highest average ratings.

Script:

`top_10_movies = LIMIT sorted_movies 10;`

`STORE top_10_movies INTO '/user/maria_dev/top_10_movies.txt';`

```
grunt> top_10_movies = LIMIT sorted_movies 10;
grunt> STORE top_10_movies INTO '/Project_Dataset/top_10_movies.txt';
```

Figure 9 Query

Output:

```
Success!

Job Stats (time in seconds):
JobId  Maps  Reduces MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime  MaxReduceTime  MinReduceTime  AvgReduceTime  MedianReduceTime  Alias  F
eature  Outputs
job_1694535283099_0078  1  1  12  12  12  12  8  8  8  8  avg_ratings_by_user,grouped_ratings_by_user,ratings  GROUP_BY
,COMBINER
job_1694535283099_0079  1  1  6  6  6  6  4  4  4  4  sorted_movies  SAMPLER
job_1694535283099_0080  1  1  5  5  5  5  4  4  4  4  sorted_movies  ORDER_BY,COMBINER
job_1694535283099_0081  1  1  4  4  4  4  3  3  3  3  sorted_movies  /Project_Dataset/top_10_movies.txt,

Input(s):
Successfully read 100000 records (1979558 bytes) from: "/Project_Dataset/u.data"

Output(s):
Successfully stored 10 records (209 bytes) in: "/Project_Dataset/top_10_movies.txt"

Counters:
Total records written : 10
Total bytes written : 209
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1694535283099_0078 -> job_1694535283099_0079,
job_1694535283099_0079 -> job_1694535283099_0080,
job_1694535283099_0080 -> job_1694535283099_0081,
job_1694535283099_0081
```


File Preview



/Project_Dataset/top_10_movies.txt/part-r-00000

849	4.869565217391305
688	4.833333333333333
507	4.724137931034483
628	4.703703703703703
928	4.6875
118	4.661971830985915
907	4.571428571428571
686	4.563380281690141
427	4.548387096774194
565	4.542857142857143

✕ Cancel

Download

3. Concluding Remark

In this project, we embarked on a comprehensive exploration of movie rating datasets, employing Apache Pig to perform various data processing tasks. The investigation involved a two-fold analysis focusing on identifying poorly rated movies and understanding user preferences through average ratings. Leveraging Pig Latin, a high-level scripting language for Apache Pig, enabled efficient data processing and analysis.

The project initiated with data loading and extraction of relevant information. The dataset was loaded, and pertinent fields such as `userID`, `movieID`, `rating`, and `ratingTime` were extracted. Subsequently, grouping and aggregation operations were executed to discern patterns in user behavior. Grouping the ratings data by `userID` facilitated the computation of average ratings for each user. The results were stored for further analysis, contributing insights into user preferences and aiding in the personalization of content recommendations. Further enriching our analysis, the project delved into the realm of data integration through joining. By merging the ratings data with metadata, which included movie titles, a more comprehensive dataset was obtained.

Sorting emerged as a crucial facet of our exploration. The sorted dataset, ordered by average ratings in descending order, unveiled the highest-rated movies. This process facilitated the identification of movies that resonated most positively with the audience, offering valuable insights into popular content.

Lastly, the project employed limiting operations to distill the wealth of information into manageable, actionable insights. By focusing on the top 10 movies with the highest average ratings, the research honed in on the most impactful findings. This selective approach is instrumental for content recommendation systems, aiding in the curation of personalized lists for users.

In conclusion, the research project demonstrated the power of Apache Pig in efficiently processing and analyzing large-scale datasets. The methodologies employed, including grouping, joining, sorting, and limiting, showcased the versatility of Pig Latin in handling diverse tasks. The outcomes of this study provide a foundation for enhancing user experiences in content recommendation systems, ensuring that users are presented with movies that align with their preferences, thereby maximizing user satisfaction. The methodologies and insights gained from this research lay a robust groundwork for future endeavors in the field of data processing and analysis.

4. Future Work

- **Sentiment Analysis:** Going beyond numerical ratings, future work could involve sentiment analysis on user reviews. Understanding the qualitative aspects of why users liked or disliked a movie could provide deeper insights into audience preferences and contribute to a more nuanced recommendation system.
- Can I also work on different sectors like Coursera reviews, Flipkart reviews and Google Play Store reviews

5. References

1. Jain, A., Bhatnagar, V., Chen, M., Mao, S., Liu, Y., Frings, S., ... Santhi, V. (2016). Movie Analytics for Effective Recommendation System using Pig with Hadoop. *International Journal of Rough Sets and Data Analysis*, 3(2), 82–100. doi:10.4018/IJRSDA.2016040106
2. Lin, J., & Dyer, C. (2010). Data-Intensive Text Processing with MapReduce. *Synthesis Lectures on Human Language Technologies*, 3(1), 1–177. DOI: 10.2200/S00218ED1V01Y201003HLT007
3. Apache Pig Documentation: The official documentation provides detailed information about Apache Pig, its functionalities, and how to use it. You can find it on the Apache Pig website: <https://pig.apache.org/docs/r0.17.0/>
4. <https://www.guru99.com/introduction-to-pig-and-hive.html>
5. [Tom White], 2009. *Hadoop: The Definitive Guide*. First Edition. O'Reilly Media.
6. [Alan Gates], 2011. *Programming Pig*. First Edition. O'Reilly Media.